

# Coding Challenge-3

## Pyspark & Spark

Name-Ajay Chaudhary

Batch-Data Engineering(Batch-1)

Execute Manipulating, Dropping, Sorting, Aggregations, Joining, GroupBy dataframes

PySpark groupBy() function is used to collect the identical data into groups on DataFrame and perform count, sum, avg, min, and max functions on the grouped data.

```
jupyter Coding_challenge_3.1 Last Checkpoint: 17 minutes ago
```

```
File Edit View Run Kernel Settings Help Trusted
```

```
Code
```

```
[1]: from pyspark.sql import SparkSession
spark=SparkSession.builder.appName("PySpark_codingChallenge").getOrCreate()
df_pyspark=spark.read.csv("/Users/ajaychaudhary/jupyter/PySpark/test.csv",header=True,inferSchema=True)
df_pyspark.show()
```

```
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
24/02/12 11:14:21 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicabl
e
```

name	department	salary
Chandu	Data Science	50000
Rashmi	IOT	75000
Rohit	Big Data	55000
Rohit	Big Data	80000
Ronit	IOT	60000
Ronit	Data Science	70000
Chandu	Data Science	45000
Krishna	Big Data	65000
Rashmi	Big Data	85000
Rohit	IOT	60000

groupBy() on Departments column of DataFrame and then the sum of salary for each department using sum() function.

```
[2]: #groupBy using sum() function
df_pyspark.groupBy("department").sum("salary").show()
```

```
[Stage 3:> (0 + 1) / 1]
```

department	sum(salary)
IOT	195000
Big Data	285000
Data Science	165000

### Aggregations-

Performing min, max, mean, avg, and count using the groupBy function.

```
[3]: df_pyspark.groupBy("department").min("salary").show()
```

department	min(salary)
IOT	60000
Big Data	55000
Data Science	45000

```
[4]: df_pyspark.groupBy("department").max("salary").show()
```

department	max(salary)
IOT	75000
Big Data	85000
Data Science	70000

```
[5]: df_pyspark.groupBy("department").avg("salary").show()
```

department	avg(salary)
IOT	65000.0
Big Data	71250.0
Data Science	55000.0

```
[6]: df_pyspark.groupBy("department").mean("salary").show()
```

department	avg(salary)
IOT	65000.0
Big Data	71250.0
Data Science	55000.0

```
[7]: df_pyspark.groupBy("department").count().show()
```

department	count
IOT	3
Big Data	4
Data Science	3

```
[8]: df_pyspark.groupBy("department").pivot("Name").sum("salary").show()
```

department	Chandu	Krishna	Rashmi	Rashmi	Rohit	Rohit	Ronit
IOT	NULL	NULL	NULL	75000	60000	NULL	60000
Big Data	NULL	65000	85000	NULL	55000	80000	NULL
Data Science	95000	NULL	NULL	NULL	NULL	NULL	70000

**Sort():** To sort a dataframe by using one or more columns, Default — ascending order

```
[9]: df_pyspark.sort("salary").show()
```

name	department	salary
Chandu	Data Science	45000
Chandu	Data Science	50000
Rohit	Big Data	55000
Ronit	IOT	60000
Rohit	IOT	60000
Krishna	Big Data	65000
Ronit	Data Science	70000
Rashmi	IOT	75000
Rohit	Big Data	80000
Rashmi	Big Data	85000

Alternatively **orderBy** can also be used to sort based on single column

```
[10]: #alternatively orderBy can be used to sort based on single column
df_pyspark.orderBy("salary").show()
```

name	department	salary
Chandu	Data Science	45000
Chandu	Data Science	50000
Rohit	Big Data	55000
Ronit	IOT	60000
Rohit	IOT	60000
Krishna	Big Data	65000
Ronit	Data Science	70000
Rashmi	IOT	75000
Rohit	Big Data	80000
Rashmi	Big Data	85000

Dropping rows based on null values

```
[11]: #for dropping
df_pyspark1=spark.read.csv("/Users/ajaychaudhary/jupyter/PySpark/null_values_test.csv",header=True,inferSchema=True)
df_pyspark1.show()
```

name	age	experience	salary
Ravi	28	5	60000
Priya	35	NULL	75000
Arun	40	12	85000
Neha	NULL	NULL	55000
Suresh	45	18	90000
Deepika	NULL	NULL	65000
Rajesh	30	8	70000
Meera	NULL	NULL	50000
Sanjay	38	15	80000
Anita	NULL	10	60000

```
[12]: #Dropping rows based on null values
df_pyspark1.na.drop().show()
```

name	age	experience	salary
Ravi	28	5	60000
Arun	40	12	85000
Suresh	45	18	90000
Rajesh	30	8	70000
Sanjay	38	15	80000

Filling missing values — Single Value

```
[13]: #filling missing values with 0
df_pyspark1.na.fill('Missing Values').show()
df_pyspark1.na.fill(0).show()
```

name	age	experience	salary
Ravi	28	5	60000
Priya	35	NULL	75000
Arun	40	12	85000
Neha	NULL	NULL	55000
Suresh	45	18	90000
Deepika	NULL	NULL	65000
Rajesh	30	8	70000
Meera	NULL	NULL	50000
Sanjay	38	15	80000
Anita	NULL	10	60000

  

name	age	experience	salary
Ravi	28	5	60000
Priya	35	0	75000
Arun	40	12	85000
Neha	0	0	55000
Suresh	45	18	90000
Deepika	0	0	65000
Rajesh	30	8	70000
Meera	0	0	50000
Sanjay	38	15	80000
Anita	0	10	60000

## Creating dataframe for performing joins.

```
[2]: from pyspark.sql import SparkSession
spark=SparkSession.builder.appName("PySpark_codingChallenge").getOrCreate()

Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
24/02/12 11:44:10 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
24/02/12 11:44:15 WARN Utils: Service 'SparkUI' could not bind on port 4040. Attempting port 4041.

[3]: #creating two dataframes for joins to be performed
emp = [(1,"John",-1,"2018","10","M",3000),(2, "Emerald",1, "2010", "20","F", 4000),(3,"Dustin",1,"2010","10","M",1000),
(4, "Nancy",2, "2005","10","F",2000),(5,"Brown",2,"2010","40","", -1),(6, "Brown", 2, "2010","50","", -1)]
empColumns = ["emp_id","name","superior_emp_id","year_joined", "emp_dept_id","gender","salary"]

empDF = spark.createDataFrame(data=emp, schema = empColumns)
empDF.printSchema()
empDF.show()

dept = [("Finance",10),("Marketing",20),("Sales",30),("IT",40)]
deptColumns = ["dept_name","dept_id"]
deptDF = spark.createDataFrame(data=dept, schema = deptColumns)
deptDF.printSchema()
deptDF.show()

root
 |-- emp_id: long (nullable = true)
 |-- name: string (nullable = true)
 |-- superior_emp_id: long (nullable = true)
 |-- year_joined: string (nullable = true)
 |-- emp_dept_id: string (nullable = true)
 |-- gender: string (nullable = true)
 |-- salary: long (nullable = true)
```

## Inner join()

Join records when key column are matched and dropped when they are not matched

```
[4]: #inner join
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"inner").show()
```

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
1	John	-1	2018	10	M	3000	Finance	10
3	Dustin	1	2010	10	M	1000	Finance	10
4	Nancy	2	2005	10	F	2000	Finance	10
2	Emerald	1	2010	20	F	4000	Marketing	20
5	Brown	2	2010	40		-1	IT	40

## Outer join()

It returns all rows from both datasets, where join expression doesn't match it returns null or respective columns.

```
[5]: #outer join
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"outer").show()
```

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
1	John	-1	2018	10	M	3000	Finance	10
3	Dustin	1	2010	10	M	1000	Finance	10
4	Nancy	2	2005	10	F	2000	Finance	10
2	Emerald	1	2010	20	F	4000	Marketing	20
NULL	NULL	NULL	NULL	NULL	NULL	NULL	Sales	30
5	Brown	2	2010	40		-1	IT	40
6	Brown	2	2010	50		-1	NULL	NULL

## Left join()

It returns all rows from left dataset regardless of match found on right dataset , when join doesn't match it assigns null for that record.

```
[6]: #left join
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"left").show()
```

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
1	John	-1	2018	10	M	3000	Finance	10
3	Dustin	1	2010	10	M	1000	Finance	10
2	Emerald	1	2010	20	F	4000	Marketing	20
4	Nancy	2	2005	10	F	2000	Finance	10
6	Brown	2	2010	50		-1	NULL	NULL
5	Brown	2	2010	40		-1	IT	40

## Right join()

It returns all rows from right dataset regardless of match found on right dataset , when join doesn't match it assigns null for that record.

```
[7]: #right join
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"right").show()
```

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
4	Nancy	2	2005	10	F	2000	Finance	10
3	Dustin	1	2010	10	M	1000	Finance	10
1	John	-1	2018	10	M	3000	Finance	10
2	Emerald	1	2010	20	F	4000	Marketing	20
NULL	NULL	NULL	NULL	NULL	NULL	NULL	Sales	30
5	Brown	2	2010	40		-1	IT	40

## Left semi join()

It returns columns from the only left dataset for the matched records in the right dataset on join expression.

```
[8]: #left semi join
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"leftsemi").show()
```

[Stage 27:===== (4 + 0) / 4]

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary
1	John	-1	2018	10	M	3000
3	Dustin	1	2010	10	M	1000
4	Nancy	2	2005	10	F	2000
2	Emerald	1	2010	20	F	4000
5	Brown	2	2010	40		-1

## Left anti join()

It returns only columns from left dataset for non- matched records.

```
[9]: #left anti join
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"leftanti").show()
```

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary
6	Brown	2	2010	50		-1